

Deletion-Based Faithfulness Metrics Do Not Discriminate Between Attribution Methods on Volumetric CT

Joana Owusu-Appiah

Erasmus Mundus MSc Medical Imaging and Applications (MAIA)

Girona, Spain • Dijon, France • Cassino, Italy

owusuappiahjoana59@gmail.com

Abstract—Attribution maps assert that the voxels they rank highest are those a model responded to. Deletion metrics test this by masking top-ranked voxels and recording the change in predicted probability. We apply the protocol to CT-FM, a SegResNet pretrained by contrastive self-supervised learning on 148,000 CT scans, fine-tuned for lung-nodule malignancy classification on volumetric chest CT, and compare three layer-wise relevance propagation composites against Grad-CAM. Two results are reported. First, CT-FM initialisation improves test AUC over an architecturally identical randomly-initialised control at every label budget, reaching 0.895 against 0.873 at full supervision; across three seeds the mean gain is positive at each budget with overlapping ranges at 10% and 25% of labels. Second, none of the four attribution methods achieved a lower deletion AUC than a random-order baseline, scoring 0.12 to 0.22 lower on AOPC. We attribute this to the masking scheme: random deletion distributes replaced voxels uniformly through the volume and moves the input away from the training distribution, whereas attribution-guided deletion removes a compact contiguous region. Under mean-value masking the baseline therefore benefits from distribution shift. Perturbation stability does separate the methods, with Grad-CAM at 0.924 and the LRP composites between -0.015 and 0.220 . A four-member deep ensemble reduces expected calibration error from 0.127 to 0.051 relative to MC dropout and attributes 53% of predictive entropy to model disagreement against 2.4% for dropout. We conclude that deletion protocols require manifold-preserving masking before they can rank attribution methods on volumetric medical data.

Index Terms—explainable AI, layer-wise relevance propagation, foundation models, self-supervised learning, uncertainty estimation, computed tomography

I. INTRODUCTION

Two claims are common in the medical imaging literature. The first is that self-supervised pretraining on large unlabelled image collections improves downstream performance. The second is that a published attribution map describes what the resulting model responded to. This paper evaluates both on a single task.

We fine-tune CT-FM [1], a SegResNet pretrained by contrastive self-supervised learning on 148,000 CT scans, for benign-versus-malignant lung-nodule classification on volumetric chest CT. We then explain it with layer-wise relevance propagation (LRP) [2] implemented in zennit [3], whose layer-type registry covers `Conv3d` and the 3D pooling operators and therefore applies to volumetric networks without modification. Grad-CAM [4] is included as a baseline.

Our contributions are:

- 1) A label-efficiency ablation against an architecturally identical from-scratch control at 10%, 25% and 100% of labels, with a seed replicate that identifies which gains survive.
- 2) A faithfulness evaluation of three LRP composites and Grad-CAM on 3D CT against a random-order deletion baseline, with an analysis of the resulting negative outcome.
- 3) A measurement of explanation stability under input perturbation, and of the correlation between attribution faithfulness and predictive uncertainty.

II. METHOD

A. Data

NoduleMNIST3D [8] comprises 3D chest CT patches centred on lung nodules drawn from LIDC-IDRI [9] and labelled benign or malignant. Splits are 1,158 training, 165 validation and 310 test volumes, with a 863:295 benign-to-malignant training imbalance. We therefore use class-weighted cross-entropy, report balanced accuracy and AUC alongside accuracy, and select models on validation AUC.

We use the 64³ release. This is not cosmetic: CT-FM downsamples by a factor of 16, so a 28³ volume is not divisible by the network stride and raises an error.

B. Model and three silent failure modes

CT-FM’s encoder produces a five-level feature pyramid; we take the bottleneck at 512 channels and 4³ resolution, apply global average pooling, LayerNorm and a linear head. Three implementation details are recorded because each affects correctness without producing an error message.

First, the published checkpoint does not load into MONAI’s SegResNetDS. Parameter names differ, and `load_state_dict(strict=False)` matches *zero of 161* tensors without raising, so training proceeds on a randomly initialised network while the code appears correct.

Second, the from-scratch control must be architecturally identical. Constructing a second SegResNet from an inferred configuration yielded 19.9M parameters against the true 87.2M; that ablation would have measured capacity and

TABLE I
LABEL-EFFICIENCY ABLATION, TEST AUC (SEED 0)

Labels	n	CT-FM	Scratch	Gain
10%	116	0.8283	0.8168	+0.0116
25%	290	0.8472	0.8124	+0.0348
100%	1,158	0.8951	0.8733	+0.0218

attributed the difference to pretraining. Our control is the loaded model with every parameter reinitialised.

Third, pooling the full SegResNet output gives the linear head two features, since that output is a 2-channel volume. With unnormalised activations, whose standard deviation is approximately 0.03, training loss remained at 6.37 with AUC 0.50.

The encoder is fine-tuned at 10^{-5} and the head at 10^{-3} . Fine-tuning 87M pretrained parameters at the head’s learning rate displaces the pretrained representation, which is the quantity under evaluation.

C. Attribution

LRP redistributes the output score backwards under conservation rules and produces a map at input resolution. We compare three composites: ε -plus-flat, ε - γ -box with box bounds matching the $[-1, 1]$ input range, and ε - $\alpha 2\beta 1$. All three are reported, since the choice of rule affects the resulting map. Grad-CAM is computed on the encoder bottleneck and trilinearly upsampled.

D. Faithfulness, stability and uncertainty

Deletion progressively replaces the highest-ranked voxels with the dataset mean and tracks the target-class probability; a faithful map produces a steep early drop and hence a low area under the curve. **Insertion** restores highest-ranked voxels onto a blurred volume. **AOPC** is the mean probability drop across deletion steps. Every map is accompanied by a random-order baseline, because absolute values depend on model and data and only the gap against random is interpretable.

Stability is the mean correlation between the map on a clean input and on a noisy copy ($\sigma = 0.05$), noise the model itself ignores.

Uncertainty is estimated by MC dropout over 20 passes, giving predictive entropy and mutual information, plus expected calibration error.

III. RESULTS

A. Self-supervised pretraining helps, with one caveat

CT-FM initialisation wins at every budget (Table I, Fig. 1), reaching 0.895 AUC and 0.830 balanced accuracy at full supervision against 0.873 and 0.788 from scratch.

Across three seeds the mean gain is +0.0198 at 10% and +0.0263 at 25% of labels, with overlapping ranges between the two arms at both budgets. One seed at 10% gave a higher score for the randomly-initialised arm. Table I should be read as a single draw, and additional replicates are required before the gains can be quoted as established.




Fig. 1. Label efficiency and the AUC gain attributable to pretraining.

TABLE II
FAITHFULNESS AND STABILITY, 64 TEST VOLUMES

Method	Del.↓	Ins.↑	AOPC—rnd	Stab.
Random baseline	0.587	—	—	—
LRP ε +flat	0.819	0.972	−0.221	0.012
LRP $\varepsilon\gamma$ box	0.822	0.972	−0.224	−0.015
LRP $\alpha 2\beta 1$	0.816	0.973	−0.219	0.220
Grad-CAM	0.712	0.974	−0.119	0.924

B. Calibration

Expected calibration error under MC dropout is 0.145 with mean predictive entropy 0.085: the model is confident and less accurate than its confidence implies, the usual direction for a small fine-tuned network on an imbalanced task.

C. No attribution method beats random deletion

A faithful map should sit *below* the random baseline on deletion AUC. All four sit well above it (Table II, Fig. 2): removing the voxels these methods rank highest damages the prediction *less* than removing random voxels.

The likely cause is the masking scheme. Random deletion distributes replaced voxels uniformly through the volume; at 5% this produces a volume unlike any CT scan, and the network’s output falls because the input is far from the training distribution. Attribution-guided deletion removes a compact contiguous region, which remains closer to a plausible scan. Under this protocol the random baseline benefits from distribution shift.

Two consequences follow. Deletion with mean-value masking does not discriminate between attribution methods on this task, and a manifold-preserving scheme such as inpainting the removed region is required before the protocol can be used for ranking. Insertion also fails to separate the methods, with all four recovering to 0.972–0.974.

D. Stability separates what deletion cannot

Under perturbation the methods diverge. Grad-CAM maps are nearly invariant (0.924) while the LRP composites give 0.012, −0.015 and 0.220. Grad-CAM is computed at 4^3 and upsampled, which constrains how much the map can change. LRP resolves to input resolution and is correspondingly more sensitive to input perturbation. Neither metric reported here establishes which behaviour is preferable.




Fig. 2. Deletion AUC against the random-order baseline, and AOPC relative to random. All four attribution methods fall on the wrong side of the baseline.




Fig. 3. Attribution on the central axial slice for three test volumes.

E. Deep ensemble against MC dropout

Each full-label run stores a tagged checkpoint, so the seed replicates used for Section III-A also serve as ensemble members. Four independently trained models were available, with individual test AUC 0.8951, 0.9093, 0.9111 and 0.8805.

The ensemble reduces expected calibration error from 0.1265 to 0.0505 at comparable AUC (Table III, Fig. 4). The two estimators differ in how they apportion the uncertainty.

TABLE III
UNCERTAINTY ESTIMATORS ON THE TEST SPLIT

Estimator	AUC	ECE	Entropy	MI share
Best single member	0.9111	—	—	—
MC dropout, 20 passes	0.8960	0.1265	0.1254	0.024
Deep ensemble, 4 members	0.9085	0.0505	0.3256	0.530




Fig. 4. Calibration error, the share of predictive entropy attributable to model disagreement, and the entropy distribution, for MC dropout against a four-member deep ensemble.

Under MC dropout, mutual information accounts for 2.4% of predictive entropy, assigning almost all reported uncertainty to ambiguity in the data. Under the ensemble it accounts for 53.0%, indicating disagreement between independently trained models that dropout sampling on a single parameter set does not expose.

F. Uncertainty against faithfulness

Spearman correlations between predictive entropy and per-sample AOPC are +0.190, +0.174 and +0.221 for the three LRP composites and -0.124 for Grad-CAM. The three LRP composites agree with each other and give weakly positive correlations; Grad-CAM gives a weakly negative one. With 64 volumes none of these correlations is statistically significant. The same question is posed for segmentation by Salahuddin et al. [6].

IV. LIMITATIONS

Table I reports a single seed, and across three seeds the ranges overlap at 10% and 25% of labels. The explanation analysis covers 64 volumes, which is sufficient to show that the deletion metric does not discriminate but not to resolve a weak correlation. The deletion baseline replaces voxels with the dataset mean; a blurred or inpainted baseline would remain nearer the data distribution and may change the conclusion of Section III-C. Inputs are 64^3 patches, not whole chest CT, so a model pretrained on complete scans may not be used as intended. The deep ensemble has four members, which is at the low end for stable uncertainty estimates, and its members share a pretrained initialisation, so their diversity arises from data ordering, subsampling and dropout only.

V. CONCLUSION

Self-supervised CT pretraining improved downstream nodule classification at every label budget, with overlapping seed ranges at the two lower budgets. The standard deletion protocol ranked all four attribution methods below a random-order baseline. We attribute this to distribution shift introduced




Fig. 5. Per-sample predictive entropy against AOPC. Correlations are weak and the three LRP composites agree with one another and disagree with Grad-CAM.

by uniform voxel masking. Faithfulness metrics developed for 2D natural images require manifold-preserving masking before they can rank attribution methods on volumetric medical data. Reporting the random control is necessary, since without it the same measurements would support a ranking of the four methods.

REPRODUCIBILITY

Code, result files and all figures are available at <https://github.com/Joana-Mansa/ct-fm-lrp-uncertainty>. The full ablation runs in approximately 14 minutes on one A100. On Volta hardware, `torch` must be pinned to 2.5.1+cu121; releases from 2.14 ship an architecture list beginning at `sm_75` and every CUDA call fails.

REFERENCES

- [1] S. Pai et al., “Vision foundation models for computed tomography,” *arXiv:2501.09001*, 2025.
- [2] S. Bach, A. Binder, G. Montavon, F. Klauschen, K.-R. Müller, and W. Samek, “On pixel-wise explanations for non-linear classifier decisions by layer-wise relevance propagation,” *PLoS ONE*, vol. 10, no. 7, 2015.
- [3] C. J. Anders, D. Neumann, W. Samek, K.-R. Müller, and S. Lapuschkin, “Software for dataset-wide XAI: from local explanations to global insights with Zennit, CoRelAy and ViRelAy,” *arXiv:2106.13200*, 2021.
- [4] R. R. Selvaraju et al., “Grad-CAM: visual explanations from deep networks via gradient-based localization,” in *ICCV*, 2017.
- [5] W. Samek, A. Binder, G. Montavon, S. Lapuschkin, and K.-R. Müller, “Evaluating the visualization of what a deep neural network has learned,” *IEEE Trans. Neural Networks and Learning Systems*, vol. 28, no. 11, 2017.
- [6] Z. Salahuddin et al., “Counterfactuals and uncertainty-based explainable paradigm for the automated detection and segmentation of renal cysts in computed tomography images: a multi-center study,” *arXiv:2408.03789*, 2024.
- [7] Z. Salahuddin, H. C. Woodruff, A. Chatterjee, and P. Lambin, “Transparency of deep neural networks for medical image analysis: a review of interpretability methods,” *Computers in Biology and Medicine*, vol. 140, 2022.
- [8] J. Yang et al., “MedMNIST v2: a large-scale lightweight benchmark for 2D and 3D biomedical image classification,” *Scientific Data*, vol. 10, 2023.
- [9] S. G. Armato III et al., “The Lung Image Database Consortium (LIDC) and Image Database Resource Initiative (IDRI): a completed reference database of lung nodules on CT scans,” *Medical Physics*, vol. 38, no. 2, 2011.
- [10] Y. Gal and Z. Ghahramani, “Dropout as a Bayesian approximation: representing model uncertainty in deep learning,” in *ICML*, 2016.